

LIVE ONLINE BOOTCAMP

Become an AI Engineer

Go from calling an API to shipping production LLM products - RAG, agents, evals and deployment. Build a portfolio of 3 real apps over 8 weekends.

For developers - Sat + Sun live - Midweek office hours - Python or JS basics is enough

8

WEEKS

3h

PER SESSION

20

MAX BATCH

Live

ONLINE

```
# build an AI agent
from rewire import Agent

agent = Agent(
    model="claude",
    tools=[search, rag],
    evals=True,
)
agent.ship() # -> prod
```

LAUNCH COHORT

~~Rs. 39,999~~

Rs .

20,000

+
GST

Enroll at rewire.school/ai-engineer



The Reality the Market Is Telling You



01

AI Engineer is the hottest role

Companies are hiring AI Engineers faster than they can find them. It is the fastest-growing role in software today.



02

Calling an API is not the job

Anyone can hit an LLM endpoint. Few can build something reliable - retrieval, agents, guardrails, and cost control.



03

Evals are the missing skill

Both OpenAI and Anthropic leaders call evals the most important new skill. Almost nobody in India teaches it. You will.

app.py

```
def answer(q):  
    ctx = rag(q)  
    return llm(  
        q, ctx  
    )  
# eval: 0.94
```

About the Bootcamp

The Rewire AI Engineer Bootcamp is built for developers who want to **build, evaluate, and ship** production LLM applications - not watch another tutorial. Over 8 weekends you move from your first API call to a deployed, evaluated AI product.

Every weekend ends with something shipped. By Demo Day you have proof - a RAG app, a tool-calling agent, and a deployed capstone with real evals - not just a certificate.

"Most devs will use AI. But you will engineer with it."

Who Is This For?

Working Developers (1-8 yrs)

Backend, frontend, or full-stack engineers who want to move into AI and stay relevant.

CS Grads Who Can Code

You know Python or JS basics and want a portfolio that gets you an AI Engineer interview.

Founders & Indie Hackers

You want to build and ship your own AI product without waiting on a team.

QA & Data Engineers Pivoting

You want to own AI quality, evals, and reliability - the part most teams are missing.

Not for absolute beginners to programming. You should be comfortable writing functions and calling an API.

Programme at a Glance

 8 WEEKS	 3 hr PER SESSION	 Live ONLINE	 RAG APP	 Agent TOOL-USE	 Demo WEEK 8
--	---	--	--	---	--

Tangible Deliverables

 RAG Application Chat-over-your-docs with retrieval, citations and grounding. Deployed live.	 Tool-Calling Agent An agent that plans and completes multi-step tasks with real tools.	 Evals + Capstone A production app with an eval suite and observability. Hiring-ready.
--	---	--



Capstone Demo Day - Week 8

Present your deployed AI product, walk through your evals, and ship your portfolio. Certificate + recording.

SAMPLE CAPSTONE PROJECTS YOU COULD SHIP

 docs-gpt <pre># RAG over wiki retrieve(q) -> chunks + cite grounded: 0.96</pre>	 support-agent <pre># auto-triage agent.run(ticket) tools: 4</pre>	 eval-bench <pre># LLM-as-judge score(out) -> pass / fail ci: green</pre>
---	--	--

Voice Assistant
realtime API

SQL Copilot
text-to-query

Research Agent
web + tools

Doc Extractor
structured out

The 8-Week Build Journey

WEEKS 1-2

LLM Foundations & Prompts

- How LLMs really work - tokens, context, sampling
- The API layer: OpenAI + Anthropic SDKs
- Structured outputs & function/tool calling
- Prompt + context engineering for production
- The SHIP framework for AI features

Output: A reliable structured-output service

WEEKS 3-4

Retrieval & RAG

- Embeddings & vector databases (pgvector, Pinecone)
- Chunking, indexing & hybrid search
- Building a RAG pipeline with citations
- Grounding & reducing hallucinations
- Deploy your RAG app over real docs

Output: Deployed RAG application

WEEKS 5-6

Agents & Tools

- Agent loops, planning & tool use
- Multi-step workflows at scale
- Memory: working vs long-term
- Guardrails, retries & error handling
- Ship a working tool-calling agent

Output: Autonomous tool-calling agent

WEEK 7

Evals & Observability

- Building eval datasets that matter
- LLM-as-judge and its limits
- Tooling: LangSmith, Braintrust, PromptFoo
- CI/CD for AI quality + the TRACE framework
- Tracing, cost & latency monitoring

Output: Eval suite wired to your app

WEEK 8: SHIP

Deploy & Capstone Demo Day

- Deployment, streaming & rate limits
- Cost optimization, caching & security
- Prompt injection defenses
- Live capstone presentation + portfolio

Ship a live, evaluated production AI app!

Proprietary Frameworks You'll Learn

SHIP

How to Build a Production AI Feature

- S**cope - define the task & success metric
- H**arness - assemble the right context
- I**nstrument - add evals before you trust it
- P**roductionize - deploy, monitor, optimize cost

TRACE

How to Evaluate & Trust an LLM

- T**est cases - build a real eval set
- R**ubric - define what good looks like
- A**utomate - LLM-as-judge where it fits
- C**ompare - score versions against each other
- E**valuate - run it in CI on every change

How the Program Works

Two weekend build sessions, a midweek office hour, and a ship-it assignment every week.

01 SATURDAY 3h Live Build Concepts taught through live coding. You build along with the instructor, not just watch.	02 SUNDAY 3h Project Lab You work on your own project with the mentor guiding and reviewing code live.	03 WED / THU Office Hours Debugging, code review, and unblocking. Recordings shared after every session.	04 SUN - FRI Ship-It Assignment Each week you ship a working piece of your portfolio. Practice that compounds.
---	--	--	--

5 Bonus Recorded Sessions

Self-paced deep dives. Released at the right moment in the cohort. Yours to keep.

01 Python for AI Crash Course <i>Before Week 1</i> async, typing, env setup - so day one is smooth.	02 Fine-tuning vs RAG <i>After Week 4</i> When to fine-tune, when to retrieve, when to do neither.
03 Prompt Injection & LLM Security <i>After Week 6</i> Attacks, defenses, and safe tool execution.	04 Cost Optimization Playbook <i>After Week 7</i> Caching, model routing, and cutting token spend 10x.
05 Landing Your First AI Engineer Role <i>After Week 8</i> How to present your portfolio, the resume bullets that work, and how to handle AI Engineer interview questions.	

All 5 bonus recordings included with every seat
 Self-paced. Released alongside the matching weekend. Yours to keep forever.

5
 VIDEOS

45m
 EACH

Tools You'll Master

LLM APIS & SDKS	BUILD & RETRIEVE	EVAL & SHIP
- Anthropic / Claude - OpenAI - Gemini	- LangChain / LlamaIndex - pgvector / Pinecone - Pydantic / FastAPI	- LangSmith / Braintrust - PromptFoo - Vercel / Railway / Modal

How We're Different

YOUTUBE / OTHER COURSES	REWIRE AI ENGINEER BOOTCAMP
- Recorded videos - Toy demos, no production - Certificate only - No evals or code review - No refund policy	48 hours LIVE instruction - Max 20 - real code review 3 shipped projects + capstone - Evals & observability included 100% refund after first session

From API-Caller to AI Engineer in 8 Weeks

BEFORE THE BOOTCAMP

Developer Who Calls an API

- Can hit an LLM endpoint but outputs are flaky
- Copies tutorials, can't adapt them to real data
- No idea how to build RAG or agents reliably
- Can't measure if the AI is actually good
- Nothing deployed to show in an interview

AFTER 8 WEEKS

Hiring-Ready AI Engineer

- + Ships reliable LLM features with structured output
- + Builds RAG pipelines and tool-calling agents
- + Writes evals and runs them in CI
- + Monitors cost, latency and quality in production
- + Has 3 deployed apps + a capstone in their portfolio

Roles You'll Be Ready For

AI Engineer

Build LLM-powered product features

Rs 18-40 LPA

LLM Application Developer

RAG, agents & integrations

Rs 15-35 LPA

AI Quality / Evals Engineer

Own reliability & evaluation

Rs 20-45 LPA

Founding Engineer / Indie

Ship your own AI product

Your own



Bonus session 05 covers your job search

Portfolio walkthrough, resume bullets that land, and how to answer AI Engineer interview questions. Salary bands are indicative of the Indian market and vary by company and experience.

After 8 Weeks, You Will Be Able To...

- ✓ Design **prompts & context** for reliable production outputs
- ✓ Ship a **tool-calling agent** that completes multi-step tasks
- ✓ Add **observability** - tracing, cost and latency
- ✓ Deploy a **streaming AI app** to production
- ✓ Build a **RAG pipeline** with retrieval and citations
- ✓ Write **eval datasets** and run LLM-as-judge
- ✓ Defend against **prompt injection** and add guardrails
- ✓ Present a **portfolio of 3 deployed AI products**

Got Questions?

Do I need to know ML or math?

No. This is AI engineering, not model training. No neural nets from scratch, no heavy math. You use models, you don't build them.

What coding background do I need?

Comfortable with Python or JavaScript basics - you can write functions and call an API. The Python crash course bonus covers the rest.

Which models and tools will I use?

Anthropic Claude and OpenAI as primary, plus open tooling for retrieval and evals. Free tiers and credits are used wherever possible.

What if I miss a session?

Every session is recorded. Mid-week office hours help you catch up. Small cohort means no one slips through the cracks.

What's the refund policy?

100% refund if it's not a fit after the first session. No questions asked.

Will this help me get hired?

You leave with 3 deployed projects + a capstone, plus a dedicated bonus on portfolio, resume and interview prep for AI Engineer roles.


```
deploy.sh

# ship your capstone
$ rewire deploy --evals
building... ok
evals: 0.94 pass
live: prod ✓
```

NEXT COHORT - LIMITED TO 20 SEATS

Your **AI Engineer** Career Starts Here

8 weekends. 3 shipped projects. One deployed capstone.

~~Rs. 39,999~~**Rs. 20,000**

+ GST

50%
OFF**Enroll at rewire.school/ai-engineer →** hello@rewire.school +91 73050 82534